

TOWARDS EFFICIENT MODELS FOR REAL-TIME DEEP NOISE SUPPRESSION

Sebastian Braun, Hannes Gamper, Chandan K.A. Reddy, Ivan Tashev

Microsoft Research, Redmond, WA, USA
sebastian.braun@microsoft.com

ABSTRACT

With recent research advancements, deep learning models are becoming attractive and powerful choices for speech enhancement in real-time applications. While state-of-the-art models can achieve outstanding results in terms of speech quality and background noise reduction, the main challenge is to obtain compact enough models, which are resource efficient during inference time. An important but often neglected aspect for data-driven methods is that results can be only convincing when tested on real-world data and evaluated with useful metrics. In this work, we investigate reasonably small recurrent and convolutional-recurrent network architectures for speech enhancement, trained on a large dataset considering also reverberation. We show interesting tradeoffs between computational complexity and the achievable speech quality, measured on real recordings using a highly accurate MOS estimator. It is shown that the achievable speech quality is a function of network complexity, and show which models have better tradeoffs.

Index Terms— speech enhancement, noise reduction, convolutional recurrent neural network, efficient neural networks

1. INTRODUCTION

Speech enhancement using neural networks has seen large attention in research in the recent years [1] and is starting to be deployed in commercial human-to-human communication applications. While the trend in research still majorly follows the trajectory of developing larger networks to further improve the performance and quality, for real-world applications following the opposite trend is of much higher interest: *How to obtain the best speech quality given a maximum computational budget?* Running current state-of-the-art noise suppression neural networks is still challenging on resource limited devices, where noise suppression is often only a small fraction among several other tasks running on the devices, such as other audio processing tasks, video, encoding, transmission, etc.

Earlier network architectures were mainly recurrent neural network (RNN) structures, which were believed promising in terms of efficiency due to its efficient temporal modeling capabilities [2–4]. While such models seem to hit a performance saturation, the use of convolutional recurrent networks (CRNs) and convolutional neural networks (CNNs) raised the performance, but resulted in development of enormously large architectures [5–8] that are impractical to run on typical edge devices like consumer laptops, mobile phones, or even less powerful devices like wearables or hearing aids. Efficient models are also obtained by building as much prior knowledge into the models as possible, rather than trying to learn well-understood blocks such as time-frequency transforms from scratch. While time-domain networks such as [9] could in theory yield superior performance than frequency-domain (FD) networks, proof of generalization on real data in reverberant environments and real recordings has

not yet been shown [10]. Therefore, we stick in this work to FD implementations.

To draw valid and general conclusions from our experiments, we train on large scale data simulating the most important aspects of real-world data such as reverberation, many different speakers, a vast amount of noise types, and varying microphone signal levels. We propose a powerful data generation and augmentation pipeline that deals with reverberant and non-reverberant speech to reduce heavy reverberation, using signal-based estimates of reverberation parameters. Results are shown on real recordings of a public dataset using a deep neural network (DNN) based MOS predictor that has shown high correlation to subjective ratings in practice.

In this work, we compare RNN with CRN architectures and show which network parts can be scaled, removed, or replaced by more efficient modules, at which gains in complexity and which loss in quality. Specifically, we investigate the influence of RNN size, type, and the use of disconnected parallel RNNs. For CRNs with a symmetric convolutional encoder/decoder structure, we investigate the convolution layers, spectral vs. spectro-temporal convolutions, and skip connections. As a result, we propose an efficient CRN structure with around 4-5 times less computational operations with similar quality than previously proposed CRNs.

2. ENHANCEMENT SYSTEM AND TRAINING OBJECTIVE

We use spectral suppression-based enhancement systems due to their robust generalization, logical interpretation and control, and easier integration with existing speech processing algorithms. The input features to the networks are log power spectra. The network predicts a real-valued, time-varying suppression gain per time-frequency bin, that is applied to the complex input spectrum, and transformed back to time-domain as shown in Fig. 1 in the upper branch. To compute a single frame, the network requires only the feature of the current frame, or when using causal convolutions, also several past frames. Therefore, the algorithmic delay of the systems depend only on the short-time Fourier transform (STFT) window size.




Fig. 1. Enhancement system and training with STFT consistency [8] and level-invariant loss [11].




Fig. 2. CRUSE network architecture with L encoder/decoder layers and a bottleneck with P parallel recurrent layers.

We train the networks enforcing STFT consistency [8] by propagating the FD output through reconstruction and another STFT to compute a FD loss as shown in Fig. 1. This preserves the flexibility of integrating the network with other FD algorithms, and offloading the STFT operations to optimized implementations. As shown in Fig. 1, each training sequence, i.e. predicted and target signals, are normalized by the active target utterance level, to ensure balanced optimization for signal-level dependent losses [11].

We train on the complex compressed mean-squared error (MSE) loss [12], blending the magnitude-only with a phase-aware term, which we found to be superior to other losses for reverberant speech enhancement [13]. The loss function per sequence is given by

$$\mathcal{L} = (1-\lambda) \sum_{k,n} \left| |S|^c - |\hat{S}|^c \right|^2 + \lambda \sum_{k,n} \left| |S|^c e^{j\varphi_S} - |\hat{S}|^c e^{j\varphi_{\hat{S}}} \right|^2, \quad (1)$$

where $c = 0.3$ is a compression factor, $\lambda = 0.3$ [13] is a weighting between complex and magnitude-based loss, and we omitted the dependency of the target speech spectral bins $S(k, n)$ on the frequency and time indices k, n for brevity.

The networks are trained in batches of 10 sequences of 10 s length using the AdamW optimizer [14], learning rate of $8 \cdot 10^{-5}$, and weight decay of 0.1. The best model is picked based on the validation metric using a heuristic optimization criterion Q using perceptual evaluation of speech quality (PESQ) [15], scale-invariant signal-to-distortion ratio (siSDR) [16] and cepstral distance (CD) [17]:

$$Q = PESQ + 0.2 \cdot siSDR - CD. \quad (2)$$

3. NETWORK ARCHITECTURES

In this section, we describe RNN and CRN architectures and modify them to improve efficiency. All models use the same features, prediction targets, loss, and training strategy described in Section 2.

3.1. NSnet2

The network proposed in [11], referred to as *NSnet2*, consists only of fully connected (FC) and gated recurrent unit (GRU) [18] layers in the format FC-GRU-GRU-FC-FC-FC. All FC layers are followed by rectified linear unit (ReLU) activations, except the last layer has sigmoid activations to predict a constrained suppression gain. The standard layer dimensions are 400 for GRUs, 600 for FC layers, i.e. 400-400-400-600-600- K , but we also investigate different configurations. The input and output dimensions are the number of frequency bins K .




Fig. 3. Skip connections by a) doubling the decoder input channels, b) addition. We found inserting 1×1 convolutions in the add skip connections useful.

3.2. CRUSE

The second network is a CRN U-Net structure derived from [7], referred to in the remainder as *Convolutional Recurrent U-net for Speech Enhancement (CRUSE)*. As shown in Fig. 2, the structure has L symmetric convolutional and deconvolutional encoder and decoder layers with kernels of size (2, 3) in time and frequency dimensions. The convolution kernels move with a stride of (1, 2), i.e. downsample the features along the frequency axis efficiently, while the number of channels C_ℓ for layer $\ell = \{1, \dots, L\}$ increase per encoder layer, and decrease mirrored in the decoder. In this work, input and output channels $C_{in} = C_{out} = 1$, but they can be extended to e.g. take complex values or multiple features as multiple channels. Convolutional layers are followed by leaky ReLU activations, while the last layer uses sigmoid. Between encoder and decoder sits a recurrent layer, which is fed with all features flattened along the channels. In [7] a stack of two long-term short-term (LSTM) layers was proposed at this stage. As will be shown in our experimental results in Section 5, replacement by a single GRU layer yields very little performance loss, but huge computational savings. A GRU saves 25% complexity compared to an LSTM layer. Two further modifications are addressed in the following two paragraphs.

Parallel RNN grouping As will be shown in Section 5, the performance of both CRUSE and NSnet2 directly scales with the bottleneck size, i.e. the width R of the central RNN layer(s). However, the complexity of RNN layers scales with R^2 , making wide RNNs computationally unattractive. Therefore, we adopt the technique proposed in [19], to group the wide fully connected RNNs into P disconnected parallel RNNs, still yielding the same forward information flow as shown in Fig. 2. We denote the number of P parallel GRUs, where $P = 1$ means the last convolutional encoder output is flattened to a single vector and fed to a single GRU, while with $P > 1$, the encoder output is reshaped to P vectors of same length, being fed through P disconnected GRUs, and being reshaped again to the number of decoder channels C_L . Another practical advantage is the possible parallel execution of the disconnected RNNs.

Skip connections As shown in Fig. 2, each convolutional encoder layer is connected to its corresponding decoder layer by a skip connection. In [19] skip connections between corresponding encoder and decoder layers have been implemented by concatenating the encoder output to the corresponding decoder input along the channel dimension as shown in Fig. 3a). This doubles the number of decoder input channels, resulting in higher complexity. More efficient skip connections are implemented by simply adding the encoder outputs to the decoder inputs, resulting in minor performance degradation. We found when adding a trainable channel-wise scaling and bias in the add-skip connections, which can be implemented as convolutions with C_ℓ channels and (1,1) kernels as in Fig. 3b), therefore being very cheap, improves the performance at negligible additional cost.

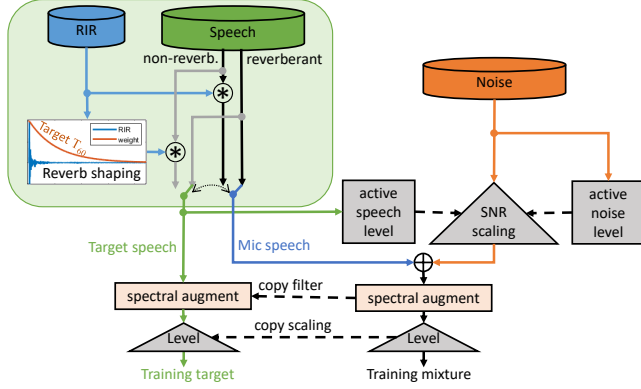


Fig. 4. Training data generation: Reverberant speech is used as, while non-reverberant speech is augmented with RIRs, and the training targets are created using shaped RIRs.

4. EXPERIMENTAL SETUP

4.1. Dataset

We use a large-scale synthetic training set and test on real recordings to ensure generalization of our results to real-world signals. The training set uses 544 h of high mean opinion score (MOS) rated speech recordings from the LibriVox corpus, 247 h noise recordings from Audioset, Freesound, internal noise recordings and 1 h of colored stationary noise. Except for the 65 h internal noise recordings, the data is available publicly as part of the 2nd DNS challenge¹. We estimated T_{60} and C_{50} for each speech file using [20, 21] and classified them as reverberant if $T_{60} > 0.22$ s and $C_{50} < 18$ dB.

Our data generation pipeline, outlined in Fig. 4, is described in the following. While already reverberant speech files are mixed with noise as is, non-reverberant speech files were augmented with acoustic impulse responses randomly drawn from a set of 7000 measured and simulated responses from several public and internal databases. 20% non-reverberant speech is not reverb augmented to represent conditions such as close-talk microphones or professionally recorded speech. To obtain natural sounding, low-reverberant, and time-aligned target speech signals, the reverberant impulse responses were shaped to a maximum decay of $T_{60}^{\max} = 0.3$ s as shown in Fig. 4. The weighting function (shown as red line in the reverb shaping block) is an exponential decay with the desired reverberation time [22], starting at the direct sound t_0 of the room impulse response (RIR)

$$w_{\text{RIR}}(t) = \begin{cases} \exp\left(-(t - t_0) \frac{6 \log(10)}{T_{60}^{\max}}\right), & t \geq t_0 \\ 1, & t < t_0 \end{cases} \quad (3)$$

To generate noisy training data, random speech and noise portions were selected to form training clips of 10 s length. Each speech and noise segment is level normalized and concatenated with other clips, if the duration was too short. The reverberant speech segments are then generated as described before, shown in the green box in Fig. 4, providing the reverberant speech and non-reverberant target speech signals. The reverberant speech and noise is mixed with a signal-to-noise ratio (SNR) drawn from a Gaussian distribution with $\mathcal{N}(5, 10)$ dB. The resulting mixture signals are re-scaled to levels distributed with $\mathcal{N}(-28, 10)$ dBFS. The speech targets are scaled

accordingly with the same factors. The optional spectral augmentation was not used here due to the large amount of raw data. Using this pipeline, we created an augmented dataset of roughly 1000 h.

For training monitoring and hyper-parameter tuning, we generated a synthetic validation set in the same way as above, using speech from the DAPS² dataset, and RIRs and noise from the QUT³ database. The final test results are shown on the public development set of the Interspeech 2020 DNS challenge [23], consisting of 400 real recordings and 300 synthetic mixtures from unseen datasets.

4.2. Evaluation metrics

Evaluating speech enhancement algorithms is a complex task, which led to the development of various objective metrics, while subjective ratings are still the gold standard. Recently, DNNs are developed to predict MOS [24, 25]. While we evaluated most of the proposed models using crowd-sourced ITU P.808 tests, we show only the predicted DNSMOS [25] for consistency of presentation and space constraints here. Nevertheless, all rankings and trends were coherent across crowd-sourced MOS, DNSMOS and intrusive objective metrics like PESQ, CD, and siSDR on the synthetic validation set.

For all models, we relate their audio quality to an estimate of the computational complexity during inference in terms of multiply-accumulate (MAC) operations. Note that we count only the operations related to applying the weights and biases, which usually contribute the major computational burden. We do not account for applying activation functions, and also omit feature extraction, STFT, and enhancement operations, which are common for all models, and are both negligible compared to the burden of the DNN models.

4.3. Algorithmic parameters

We use a sampling frequency of 16 kHz, an STFT with 50% overlapping squareroot-Hann windows of 20 ms length, and a corresponding FFT size of 320 points. The inputs to the networks are 161-dimensional log power spectra. We parameterize *NSnet2* models denoted by *NSnet2-R*, where *R* denotes the number of GRU nodes per layer. We parameterize *CRUSE* with different encoder/decoder sizes, starting always with $C_1 = 16$ channels, and doubling the channels each layer. *CRUSE* models are denoted by *CRUSEL-C_L-NxRNNP*, where *L* are the number of encoder/decoder layers, the last encoder layer filters C_L can vary to scale the RNN layer width, *N* are the number of RNN layers, and *P* are the number of parallel RNNs. For example, *CRUSE4-120-1xGRU4* has 4 encoder/decoder layers with filters 16-32-64-120, and 1 layer of 4 parallel GRUs. Convolution kernels are always (2,3), unless denoted explicitly as 1D convolutions with (1,3) kernels operating only across frequency.

5. RESULTS

Figure 5 shows the tradeoff between computational complexity in terms of MACs vs. the predicted overall audio quality using DNSMOS [25] for several *NSnet2* and *CRUSE* models, and other baselines. In Fig. 5 all *CRUSE* models use add skips without 1×1 conv blocks unless denoted explicitly. The first baseline is a classic noise suppressor (*classicNS*) exploiting only stationarity of speech [26]. While this method non-surprisingly achieves only minor MOS improvement of around 0.12 on the test set including many highly non-stationary noise types, the achievement relative to its computational

²<https://ccrma.stanford.edu/gautham/Site/daps.html>

³<https://research.qut.edu.au/saivt/databases/qut-noise-databases-and-protocols/>

¹<https://github.com/microsoft/DNS-Challenge>




Fig. 5. MOS improvement vs. computational complexity (MACs).

burden in the order of a few 1000 MACs is a tiny fraction of even our smallest models. For direct comparison of the following DNN-based baselines, we only took the architectures, but trained them using the same features, prediction targets, loss, and training procedure as for all other networks in this paper: i) *NSnet* [4] yields similar MOS as the *classic NS*. ii) The fully convolutional architecture proposed in [27] underperforms in this task, which we mainly suspect to the absence long-term temporal modeling as it uses only 8 frames temporal context. iii) The *CRN-LSTM9* architecture proposed in [28] yields significant MOS improvement, but is computationally extremely inefficient with 283M MACs due to large CNN filter kernels and wide recurrent layers.

NSnet2, shown with different RNN sizes denoted as *NSnet2-400* and *NSnet2-500*, performs better than *NSnet* and its quality can be improved by increasing the RNN size to 500 units, which also scales up the MACs. The *CRUSE5* models gain largely in efficiency by replacing the 2 LSTM layers (\times) with 2 GRUs (\times), and further by going to only a single GRU (\times). We can observe how moving from fully connected GRUs ($P = 1$) to $P = 4$ parallel GRUs for different RNN sizes ($\bigcirc \rightarrow \square$ and $\bigcirc \rightarrow \square$) reduces complexity with very little performance loss. The 2D convolutional feature extraction of *CRUSE* is very useful and efficient as moving to 1D convolutions using kernels of (1,3) deteriorates the tradeoff significantly for *CRUSE5-256-1GRU* ($\times \rightarrow \triangle$) and *CRUSE4-120-1GRU* ($\bigcirc \rightarrow \triangle$). This seems to contradict findings in [19], where no performance degradation is claimed by using 1D convolutions on *non-reverberant datasets*, which highlights the importance of using appropriate data.

Overall, Figure 5 reveals a few very interesting trends: i) The fully recurrent *NSnet* models have generally lower complexity, but also lower quality than the *CRUSE* models. We can conclude that the convolutional encoder-decoder structures of *CRUSE* are very helpful. Especially using also temporal encoder/decoder layers boosts efficiency. ii) We can observe a surprisingly linear correlation between MACs and MOS. Consequently for the tested models, the average achieved quality is a monotonically increasing function of the computational effort. The MAC-MOS linear relation is even more clear for models within the same class, e.g. the *NSnet* models, or *CRUSE* models with different RNN widths. The most efficient models, i.e. the proposed *CRUSE* models with parallel GRUs and optimized skip connections (\triangleright) break out of the linear trend, pushing towards the desired upper left corner.

Fig. 6 illustrates scalable performance depending on the RNN width. The blue and red lines show *NSnet2-R* and *CRUSE4-C₄-1xGRU1* models, where the width of the RNN was scaled by changing the RNN width R for *NSnet2*, or changing the last encoder layer's filters C_4 and therefore also the RNN width for *CRUSE4*, respectively. We can observe a clear monotonic relationship between




Fig. 6. Tradeoff by changing RNN width for *NSnet2* and *CRUSE*, and parallel RNN grouping. Colored numbers denote the variable parameter per model, R , C_4 , and P .

Table 1. Effect of skip connection type for *CRUSE4-128-1GRU4*.

model	MACs (M)	Δ MOS
no skips	4.3	0.32
add skips	4.3	0.35
add conv 1×1 skips	4.3	0.38
concat skips	4.8	0.38

RNN throughput or memory and the obtained quality. Obviously, the complexity vs. quality tradeoff deteriorates for too large models. It is a useful property that the performance of these networks can be scaled given a certain computational budget. The yellow line in Fig. 6 shows the effect of grouping the fully connected RNN for a *CRUSE4-128-1xGRUP* model into P disconnected RNNs while keeping the RNN feedforward flow fixed. Significantly better tradeoffs can be achieved with $P = 2$ and $P = 4$.

Table 1 shows further ablations on the skip connections using the *CRUSE4-128-1xGRU4P* architecture. Addition skips are better than no skip connections. While concat skips improve the MOS, the same MOS can be achieved by inserting cheap 1×1 convolutions as shown in Fig. 3 by the orange block.

While the execution time of the models is subject to optimization for the targeted hardware platform, we provide a sense of relating the MACs to the actual execution time of the efficient *CRUSE4-128-1xGRU4* model measured on a Intel[®] Core[™]i7 QuadCore at 3.5 GHz: Without further optimization, the ONNX model processes one audio frame on average in 0.3 ms, resulting in a reasonable CPU utilization of less than 3% within the hop size budget of 10 ms. *NSnet2-500* runs in 0.15 ms.

6. CONCLUSIONS

We proposed a flexible and scalable CRN for speech enhancement, trained on large data and tested on real recordings. We show that using simple spectral suppression based networks of comparatively small size can achieve substantial quality improvement when trained on a representative dataset with a suitable loss taking the time domain reconstruction into account. We show that the obtained speech quality is a function of network size, especially depending on the recurrent layer width. We show gains on the speech quality vs. computational complexity tradeoff by modified skip connections and a disconnected parallel RNN structure. While the proposed models use only a fraction of the computational budget of standard CPUs in real-time, the quality gain per computational burden compared to a traditional speech enhancement method is still less efficient and can hopefully be further improved in the future.

7. REFERENCES

- [1] D. Wang and J. Chen, "Supervised speech separation based on deep learning: An overview," *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, vol. 26, no. 10, pp. 1702–1726, Oct 2018.
- [2] F. Weninger, H. Erdogan, S. Watanabe, E. Vincent, J. Le Roux, J. R. Hershey, and B. Schuller, "Speech enhancement with LSTM recurrent neural networks and its application to noise-robust ASR," in *Proc. Latent Variable Analysis and Signal Separation*, 2015, pp. 91–99.
- [3] D. S. Williamson and D. Wang, "Time-frequency masking in the complex domain for speech dereverberation and denoising," *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, vol. 25, no. 7, pp. 1492–1501, July 2017.
- [4] R. Xia, S. Braun, C. Reddy, H. Dubey, R. Cutler, and I. Tashev, "Weighted speech distortion losses for neural-network-based real-time speech enhancement," in *Proc. IEEE Intl. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, 2020.
- [5] G. Wichern and A. Lukin, "Low-latency approximation of bidirectional recurrent networks for speech denoising," in *Proc. IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, Oct 2017, pp. 66–70.
- [6] M. Strake, B. Defraene, K. Fluyt, W. Tirry, and T. Fingscheidt, "Separated noise suppression and speech restoration: LSTM-based speech enhancement in two stages," in *Proc. IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, Oct 2019, pp. 239–243.
- [7] K. Tan and D. Wang, "A convolutional recurrent neural network for real-time speech enhancement," in *Proc. Interspeech*, 2018, pp. 3229–3233.
- [8] S. Wisdom, J. R. Hershey, K. Wilson, J. Thorpe, M. Chinen, B. Patton, and R. A. Saurous, "Differentiable consistency constraints for improved deep speech enhancement," in *Proc. IEEE Intl. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, May 2019, pp. 900–904.
- [9] Y. Luo and N. Mesgarani, "Conv-TasNet: Surpassing ideal time–frequency magnitude masking for speech separation," *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, vol. 27, no. 8, pp. 1256–1266, Aug 2019.
- [10] M. Maciejewski, G. Wichern, E. McQuinn, and J. L. Roux, "WHAMR!: Noisy and reverberant single-channel speech separation," in *Proc. IEEE Intl. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, 2020, pp. 696–700.
- [11] S. Braun and I. Tashev, "Data augmentation and loss normalization for deep noise suppression," in *Proc. Speech and Computers*, 2020.
- [12] A. Ephrat, I. Mosseri, O. Lang, T. Dekel, K. Wilson, A. Hassidim, W. T. Freeman, and M. Rubinstein, "Looking to listen at the cocktail party: A speaker-independent audio-visual model for speech separation," *ACM Trans. Graph.*, vol. 37, no. 4, July 2018.
- [13] S. Braun and I. Tashev, "A consolidated view of loss functions for supervised deep learning-based speech enhancement," arXiv:2009.12286, 2020.
- [14] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *International Conference on Learning Representations*, 2019.
- [15] ITU-T, "Perceptual evaluation of speech quality (PESQ), an objective method for end-to-end speech quality assessment of narrowband telephone networks and speech codecs," Feb. 2001.
- [16] J. L. Roux, S. Wisdom, H. Erdogan, and J. R. Hershey, "SDR - half-baked or well done?," in *Proc. IEEE Intl. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, May 2019, pp. 626–630.
- [17] N. Kitawaki, H. Nagabuchi, and K. Itoh, "Objective quality evaluation for low bit-rate speech coding systems," *IEEE J. Sel. Areas Commun.*, vol. 6, no. 2, pp. 262–273, 1988.
- [18] K. Cho, B. Van Merriënboer, D. Bahdanau, , and Y. Bengio, "On the properties of neural machine translation: Encoder-decoder approaches," in *Proc. 8th Workshop on Syntax, Semantics and Structure in Statistical Translation (SSST-8)*, 2014.
- [19] K. Tan and D. Wang, "Learning complex spectral mapping with gated convolutional recurrent networks for monaural speech enhancement," vol. 28, pp. 380–390, 2020.
- [20] H. Gamper and I. J. Tashev, "Blind reverberation time estimation using a convolutional neural network," in *Proc. Intl. Workshop Acoust. Signal Enhancement (IWAENC)*, 2018, pp. 136–140.
- [21] H. Gamper, "Blind C50 estimation from single-channel speech using a convolutional neural network," in *Intl. Workshop on Multimedia Signal Processing (MMSP)*.
- [22] J. D. Polack, *La transmission de l'énergie sonore dans les salles*, Ph.D. thesis, Université du Maine, Le Mans, France, 1988.
- [23] C. K. A. Reddy, E. Beyrami, H. Dubey, Gopal V., R. Cheng, R. Cutler, S. Matushevych, R. Aichner, A. Aazami, S. Braun, Rana P., S. Srinivasan, and J. Gehrke, "The interspeech 2020 deep noise suppression challenge: Datasets, subjective speech quality and testing framework," in *to appear in Proc. Interspeech 2020*.
- [24] A. R. Avila, H. Gamper, C. Reddy, R. Cutler, I. Tashev, and J. Gehrke, "Non-intrusive speech quality assessment using neural networks," in *Proc. IEEE Intl. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, 2019, pp. 631–635.
- [25] C. K. A. Reddy, H. Dubey, V. Gopal, R. Cutler, S. Braun, H. Gamper, R. Aichner, and S. Srinivasan, "ICASSP 2021 deep noise suppression challenge," in *arXiv:2009.06122 [eess.AS]*.
- [26] Ivan Tashev, *Sound Capture and Processing: Practical Approaches*, Wiley, July 2009.
- [27] S. R. Park and J. W. Lee, "A fully convolutional neural network for speech enhancement," in *Proc. Interspeech*, 2017.
- [28] M. Strake, B. Defraene, K. Fluyt, W. Tirry, and T. Fingscheidt, "Fully convolutional recurrent networks for speech enhancement," in *Proc. IEEE Intl. Conf. on Acoustics, Speech and Signal Processing (ICASSP)*, 2020, pp. 6674–6678.